

Train and val set sizes

Three datasets used in training. Counts are images actually loaded (real + full-synthetic only). SID drops tampered images; WildFake drops the empty wukong real source.

SID_Set 143,533 train 119,590 · val 23,943	WildFake 3,305,022 train 3,238,953 · val 66,069	CIFAKE 110,000 train 100,000 · val 10,000
---	--	--

Train vs val

Dataset	Train	Val	Train real	Train fake	Val real	Val fake
SID_Set	119,590	23,943	59,802	59,788	11,984	11,959
WildFake	3,238,953	66,069	732,794	2,506,159	14,956	51,113
CIFAKE	100,000	10,000	50,000	50,000	5,000	5,000

SID and CIFAKE are nearly 1:1 real/fake. WildFake is generator-heavy (~3.4 fake per real); training uses type-balanced batches, not this raw ratio.

Scale (log bars)

■ Train ■ Val

SID_Set



WildFake



CIFAKE



Bar length is log10(count) relative to WildFake train. SID is the main pretrain set; WildFake is the cross-generator finetune; CIFAKE is the 32x32 pipeline set (train 100k / val 10k).